Hybrid Attention in Kimi K3

Deep Dive into Kimi K3 — Chapter 5

Yin Yang

Foundation Books Project

The one idea of this chapter

K3 combines compressed recurrent memory with occasional token-addressable access; the layer *schedule* decides where each capability — and each cost — appears.

The one idea of this chapter

K3 combines compressed recurrent memory with occasional token-addressable access; the layer *schedule* decides where each capability — and each cost — appears.

Goal: type the two memory contracts, build the ordered 3:1 schedule, show what global layers make addressable, then account for memory, work, and the reported evidence — each in its own unit.

Roadmap

Two contracts, one schedule

Access, position, and memory

Cost and evidence

Two contracts, one schedule

Fixed state at most depths, records at a few



toy $(K^3G)^2$; K3 $(K^3G)^{23}G = 69K + 24G$

- Every layer maps $\mathbb{R}^d \to \mathbb{R}^d$ and owns its own state or ordered cache.

Fixed state at most depths, records at a few



toy $(K^3G)^2$; K3 $(K^3G)^{23}G = 69K + 24G$

- Every layer maps $\mathbb{R}^d \to \mathbb{R}^d$ and owns its own state or ordered cache.
- A global layer appends one record per position, then reads all legal records; addressability is a storage fact, not retrieval.

Fixed state at most depths, records at a few



toy $(K^3G)^2$; K3 $(K^3G)^{23}G = 69K + 24G$

- Every layer maps $\mathbb{R}^d \to \mathbb{R}^d$ and owns its own state or ordered cache.
- A global layer appends one record per position, then reads all legal records; addressability is a storage fact, not retrieval.

Check order before counts: a misplaced G keeps the totals and changes the model.

Access, position, and memory

What global layers add, and what they cost

- KDA compresses its own history but emits a representation at every position; the next global layer stores each one.

What global layers add, and what they cost

- KDA compresses its own history but emits a representation at every position; the next global layer stores each one.
- NoPE removes only the global position transform; causal permission and order-sensitive incoming content remain.

What global layers add, and what they cost

- KDA compresses its own history but emits a representation at every position; the next global layer stores each one.
- NoPE removes only the global position transform; causal permission and order-sensitive incoming content remain.
- Memory is affine: toy $264 + 4t$ vs. all-global $16t$ — they cross at $t = 22$; whole-K3 slope fraction $8/31$.

What global layers add, and what they cost

- KDA compresses its own history but emits a representation at every position; the next global layer stores each one.
- NoPE removes only the global position transform; causal permission and order-sensitive incoming content remain.
- Memory is affine: toy $264 + 4t$ vs. all-global $16t$ — they cross at $t = 22$; whole-K3 slope fraction $8/31$.

A smaller slope need not mean a smaller total at a short prefix.

Cost and evidence

Phase-specific costs, setup-specific results

- Prefill keeps a quadratic term at global depths; decode pays a prefix scan only at global layers.

Phase-specific costs, setup-specific results

- Prefill keeps a quadratic term at global depths; decode pays a prefix scan only at global layers.
- Ratio ablation: 3:1 has the lowest validation PPL (5.65) among five tested settings.

Phase-specific costs, setup-specific results

- Prefill keeps a quadratic term at global depths; decode pays a prefix scan only at global layers.
- Ratio ablation: 3:1 has the lowest validation PPL (5.65) among five tested settings.
- Reported vs. MLA: 2.3–2.9× faster prefill, 2.3× lower batch-1 TPOT at 1M; matched-model quality wins are not uniform.

Phase-specific costs, setup-specific results

- Prefill keeps a quadratic term at global depths; decode pays a prefix scan only at global layers.
- Ratio ablation: 3:1 has the lowest validation PPL (5.65) among five tested settings.
- Reported vs. MLA: 2.3–2.9× faster prefill, 2.3× lower batch-1 TPOT at 1M; matched-model quality wins are not uniform.

Next: Chapter 6 opens the gated MLA record that the global layers keep.