Native Multimodal Pretraining

Deep Dive into Kimi K3 — Chapter 16

Yin Yang

Foundation Books Project

The one idea of this chapter

A visual pathway is not a training recipe:
targets, ledgers, and gates must make visual evidence matter.

The one idea of this chapter

A visual pathway is not a training recipe:
targets, ledgers, and gates must make visual evidence matter.

Goal: type the training record, separate conditioning from supervision, name every mixture denominator, schedule media and stages, and keep acquisition, retention, and transfer as separate verdicts.

Roadmap

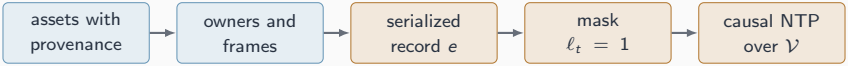
Records and supervision

Ledgers, media, and stages

Stability, transfer, and evidence

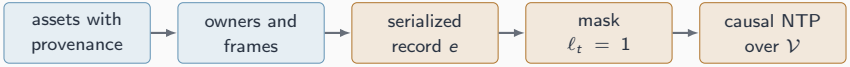
Records and supervision

From assets to causal targets



- One studio scene yields five records — text, caption, grounding, code–render, video — each with a counterfactual that must change its target.

From assets to causal targets



- One studio scene yields five records — text, caption, grounding, code–render, video — each with a counterfactual that must change its target.
- Visual positions condition; vocabulary positions are supervised. Every active target's evidence precedes it, and geometry follows the ordered transform.

Ledgers, media, and stages

One mixture, many denominators

- Equal text/visual records: token shares $1/4, 3/4$; target shares $2/3, 1/3$; compute shares $1/10, 9/10$; global visual share $1/2$. [DERIVED: two-family fixture]

One mixture, many denominators

- Equal text/visual records: token shares $1/4, 3/4$; target shares $2/3, 1/3$; compute shares $1/10, 9/10$; global visual share $1/2$. [DERIVED: two-family fixture]
- Resolution and duration are curricula: 16, 64, or 16 backbone positions for low, high, and four-frame media, gated by target-support survival.

One mixture, many denominators

- Equal text/visual records: token shares $1/4, 3/4$; target shares $2/3, 1/3$; compute shares $1/10, 9/10$; global visual share $1/2$. [DERIVED: two-family fixture]
- Resolution and duration are curricula: 16, 64, or 16 backbone positions for low, high, and four-frame media, gated by target-support survival.
- A stage is a nine-field tuple; K3 trains vision, projector, and backbone jointly from the initial state. [REPORTED: K3 Secs. 2.4, 3.3]

Stability, transfer, and evidence

Shocks, retention, and transfer

- Record the changed fields before reading a curve; diagnose from record and geometry up to optimization and capability.

Shocks, retention, and transfer

- Record the changed fields before reading a curve; diagnose from record and geometry up to optimization and capability.
- Acquisition and retention are separate verdicts; co-improvement becomes transfer only under a matched intervention.

Shocks, retention, and transfer

- Record the changed fields before reading a curve; diagnose from record and geometry up to optimization and capability.
- Acquisition and retention are separate verdicts; co-improvement becomes transfer only under a matched intervention.
- K3's masks, mixture, media laws, budgets, and gates remain undisclosed; studio numbers are teaching values. [OPEN: K3 recipe]

Shocks, retention, and transfer

- Record the changed fields before reading a curve; diagnose from record and geometry up to optimization and capability.
- Acquisition and retention are separate verdicts; co-improvement becomes transfer only under a matched intervention.
- K3's masks, mixture, media laws, budgets, and gates remain undisclosed; studio numbers are teaching values. [OPEN: K3 recipe]

Validate in dependency order. Next: Chapter 17, Supervised Cold Start for Agentic RL.